

# Introduction to Differential Privacy

Björn Bebensee  
bebensee@bi.snu.ac.kr

Seoul National University

November 14, 2019

# Contents

1. Motivation
2. Differential Privacy
  - Definition
  - Achieving Differential Privacy
  - Problems
3. Local Differential Privacy
  - Randomized response
  - Definition
  - Challenges
4. Key problems in LDP
  - Overview
  - Frequency oracles
  - Heavy hitter identification
  - Private spatial data collection
5. DP in practice

# Motivation

Imagine we have a database of users. We want to publish some anonymous statistics about the users.

# Motivation

Imagine we have a database of users. We want to publish some anonymous statistics about the users.

**This is a hard problem!**

Famously researchers uniquely identified 99% of users in the “anonymized” Netflix Prize dataset in 2007.

# Motivation

**Given:** a database of users

**Compute:** statistics over all users

**Constraint:** protect the information of each single user

**Adversary:** Potentially malicious actor who either

- is able to query the *statistical database* (*interactive setting*)
- only has access to the published statistics (*noninteractive setting*)

# Differential Privacy

Differential privacy is a framework of statistical techniques.  
It is used to

- (1) compute statistical queries on user inputs
- (2) protect each individuals' privacy while doing so

Allows a trade-off between utility and user privacy.

# Differential Privacy

## Definition

A randomized function  $\mathcal{K}$  gives  $\epsilon$ -Differential Privacy if for all data sets  $D_1$  and  $D_2$  differing on at most one element, and all  $S \subseteq \text{Range}(\mathcal{K})$ ,

$$\frac{\Pr[\mathcal{K}(D_1) \in S]}{\Pr[\mathcal{K}(D_2) \in S]} \leq e^\epsilon$$

An algorithm is  $\epsilon$ -DP if it does not depend on any single entry in the dataset but rather (probably) gives the same output even if you remove any single entry.

# Differential Privacy

How can we achieve  $\epsilon$ -Differential Privacy?



# Differential Privacy

**Idea:** Inject random noise to protect the privacy of individuals.

Example: What are the number of users with a bad credit rating?

Ground truth  $N = 21$ ,  $X$  random variable with Laplace distribution. Return  $N + X$ .

# Differential Privacy

**Idea:** Inject random noise to protect the privacy of individuals.

Example: What are the number of users with a bad credit rating?

Ground truth  $N = 21$ ,  $X$  random variable with Laplace distribution. Return  $N + X$ .

## Problem

Privacy losses accumulate. An adversary can estimate the ground truth given enough queries. More queries correspond to laxer privacy guarantees.

# Differential Privacy

To address this define a function's sensitivity as follows:

## Definition

For a function  $f : D \rightarrow R^k$ , the sensitivity of  $f$  is

$$\Delta f = \max_{D_1, D_2} \|f(D_1) - f(D_2)\|_1$$

for all datasets  $D_1, D_2$  differing in at most one element.

Then, for a query function  $f$ , return

$$f(x) + \text{Lap}(\Delta f / \epsilon)$$

with *privacy loss*  $\epsilon$ , i.e. the variance depends on the sensitivity and the privacy loss.

# Problems

However, there are some problems with (centralized) DP:

1. Still requires trust in a central authority!
2. Distributed setting: inputs are connected to identifiers such as IP address in logs etc.

# Problems

However, there are some problems with (centralized) DP:

1. Still requires trust in a central authority!
2. Distributed setting: inputs are connected to identifiers such as IP address in logs etc.

## Goal

- Compute user statistics in a distributed setting, w/o trust in central authority
- Efficient computation, low communication cost (important on mobile devices)

# Randomized response

Idea for the distributed setting: randomized response

**Survey technique** introduced in 1965 to get accurate statistics on sensitive topics.

Example: Have you used drugs this month?

# Randomized response

1. Participants toss a coin
2. Answer truthfully if coin comes up heads
3. If tails, participant tosses a 2nd coin:
  - “Yes” if tails
  - “No” if heads

True answer if first coin heads or 50% chance on second coin.  
Answer truthfully  $\sim 75\%$  of the time.

# Randomized response: analysis

Given positive answers  $X$ , true number of “Yes” answers:

$$Y = 2(X - 0.25)$$

**Result:** plausible deniability for participants, accurate statistics



# Local Differential Privacy

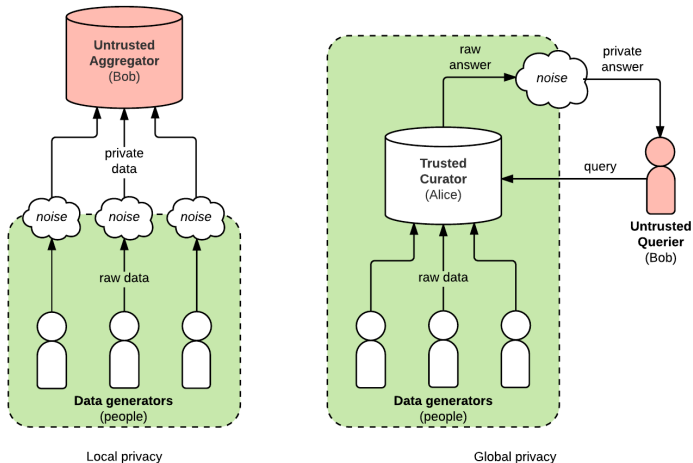


Figure: Local differential privacy vs. differential privacy

# Local Differential Privacy

## Definition

We say that an algorithm  $\pi$  satisfies  $\epsilon$ -Local Differential Privacy where  $\epsilon > 0$  if and only if for any input  $v$  and  $v'$

$$\forall y \in \text{Range}(\pi) : \frac{\Pr[\pi(v) = y]}{\Pr[\pi(v') = y]} \leq e^\epsilon$$

where  $\text{Range}(\pi)$  denotes every output of the algorithm  $\pi$ .

For the randomized response example: survey participants are given  $\epsilon$ -LDP guarantee with  $\epsilon = \ln(0.75/(1 - 0.75)) = \ln(3)$

# Challenges in the local model

## Challenges

It is much harder to construct protocols for the local model:

- more difficult to maintain low error-bound
- maintain low communication cost (ideally only a few bits)

# Challenges in the local model

## Challenges

It is much harder to construct protocols for the local model:

- more difficult to maintain low error-bound
- maintain low communication cost (ideally only a few bits)

**centralized model:** add Laplace noise of magnitude  $O(1/\epsilon)$ , independent of number of participants

**local model:** lower-bound  $\Omega(\sqrt{n}/\epsilon)$ , dependent on number of participants

# Key problems in LDP

There are a number of different problems in LDP. We will look at a few of them:

- Frequency estimation
- Heavy hitter identification
- Private spatial data collection

# Frequency Oracles

## Definition

Given a domain  $\mathcal{D}$  a *frequency oracle* (FO) is a protocol which estimates the frequency of an element  $d \in \mathcal{D}$ .

Core problem of LDP: locally private frequency estimation

# FO framework

Define *pure* FO protocols by a composition of **encoding** and **perturbation** and **support**.

**Encoding:** Encoding of each value in the domain  $\mathcal{D}$

**Perturbation:** Random perturbation of encoded values

**Support:** Mapping of each output  $y$  to the set of inputs that support the output value  $y$

The true frequency  $c(i)$  of a value  $i$  can then be estimated given the encoding, perturbation and support.

# Direct Encoding

Generalization of randomized response with coin flip.

Encode( $v$ ) =  $v$ . Perturb with:

$$\Pr[\text{Perturb}(x) = i] = \begin{cases} p = \frac{e^\epsilon}{e^\epsilon + |\mathcal{D}| - 1}, & \text{if } i = x \\ q = \frac{1-p}{|\mathcal{D}| - 1} = \frac{1}{e^\epsilon + |\mathcal{D}| - 1}, & \text{if } i \neq x \end{cases}$$

The support function is  $\text{Support}_{\text{DE}}(i) = \{i\}$ .



# Unary Encoding

Value  $v$  is encoded as a bit vector where only the  $v$ -th position equals 1 and all other positions equal 0. Given probabilities  $p, q$  the perturbed output  $B'$  is computed as follows:

$$\Pr[B'[i] = 1] = \begin{cases} p, & \text{if } B[i] = 1 \\ q, & \text{if } B[i] = 0 \end{cases}$$

with optimal parameters  $p = \frac{1}{2}$  and  $q = \frac{1}{e^\epsilon + 1}$  and the support function  $\text{Support}_{\text{UE}}(B) = \{i \mid B[i] = 1\}$ .

# Heavy hitter identification

## Goal

Estimate the frequency of the most common domain elements (*heavy hitters*) while guaranteeing  $\epsilon$ -LDP.

For small domains: simply estimate frequency of all domain elements using FO protocol

Computationally infeasible for large domains.

# Heavy hitter identification

## Definition

Set of  $n$  users each holding an input  $x_i \in \mathcal{D}$ , *distributed database*  $S = (x_1, \dots, x_n)$  consisting of all users' inputs. A domain element  $x \in \mathcal{D}$  is  $\Delta$ -heavy if its multiplicity in  $S$  is at least  $\Delta$ .

## Goal

Find all  $\Delta$ -heavy elements (i.e. *heavy hitters*) for  $\Delta$  as small as possible.

Informally:  $\Delta$ -heavy if there are at least  $\Delta$  users who hold the input  $x$ . Then  $\Delta$  is also referred to as the protocol's *error*.

# Heavy hitter identification

## Solution

Use efficient FO protocol, minimize queries necessary to identify the most frequent items.

One approach: use binary prefix tree, prune all subtrees that cannot be prefixes of heavy hitters

# Private spatial data collection

Services such as *Google Maps* and *Waze* benefit from user location data to identify popular locations and to create traffic congestion maps.

# Private spatial data collection

## Goal

We want to maintain users' privacy and be able to give privacy guarantees while collecting useful spatial data.

Unfortunately: domain too big to obtain accurate results while maintaining  $\epsilon$ -LDP

Chen et al. introduce *personalized local differential privacy*

# Private spatial data collection

## Definition

Given the personalized privacy specification  $(\tau, \epsilon)$  of a user  $u$ , a randomized algorithm  $\pi$  satisfies  $(\tau, \epsilon)$ -personalized local differential privacy for  $u$ , if for two locations  $l, l' \in \tau$  and any  $O \subseteq \text{Range}(A)$ ,

$$\frac{\Pr[\pi(l) \in O]}{\Pr[\pi(l') \in O]} \leq e^\epsilon$$

where the probability space is over the coin flips of  $\pi$ .

$\tau$  determines user's *safe region* which they do not mind revealing (i.e. “I am in New York state” but not more fine-grained than that)

PLDP is a generalized version of  $\epsilon$ -LDP as  $\tau = \mathcal{L}$  for a location universe  $\mathcal{L}$  implies regular  $\epsilon$ -LDP.

# Private spatial data collection

## Solution

- Each user chooses safe region  $\tau$
- Perturb reported locations using local randomizer which guarantees  $(\tau, \epsilon)$ -PLDP.
- Partition users into clusters with same safe region to minimize error
- Estimate user counts for all locations by combining the estimates from the clusters accordingly



# Practical deployments of DP

Local Differential Privacy is still a very active area of research.

Where is (Local) Differential Privacy in use today?

# Practical deployments of DP

In Google Chrome (RAPPOR):

- collecting user statistics
- gaining insights in common settings, how Chrome is used
- find common homepages

Basic RAPPOR uses a variant of Unary Encoding as its FO protocol

# Practical deployments of DP

On iOS and MacOS:

- for general usage statistics
- to find popular emojis
- to find new trending words that are not in the dictionary yet
- etc.

**Client side:** randomized response on bit vector encodings

**Aggregation:** uses *count-mean sketch with Hadamard transform*, requires only a single bit (!) to be transmitted by the client

# Practical deployments of DP

**Criticism:** lack of transparency in Apple's implementation, weak privacy guarantees, high privacy loss of  $\epsilon = 16$  per day

However: important first steps to adoption of (local) differential privacy have been taken, privacy guarantees can be given at scale

Thank you for your attention.