

Bridging the Gap between Training and Inference for Neural Machine Translation

Wen Zhang, Yang Feng, Fandong Meng, Di You, Qun Liu

Björn Bebensee

bebensee@bi.snu.ac.kr



Biointelligence Laboratory

November 14, 2019

Overview

A brief introduction to NMT

A brief introduction to NMT

A brief introduction to NMT

A brief introduction to NMT

A brief introduction to NMT

Notation:

Source sequence $x = \{x_1, \dots, x_{|x|}\}$

Word embeddings e_{x_i} for each $x_i \in x$

Observed translation $y^* = \{y_1^*, \dots, y_{|y^*|}^*\}$

A brief introduction to NMT

Encoder: bidirectional Gated Recurrent Unit (GRU), obtain hidden states $h_i = [\vec{h}_i; \overset{\leftarrow}{h}_i]$ where

$$\vec{h}_i = \mathbf{GRU}(e_{x_i}, \vec{h}_{i-1}) \quad (1)$$

$$\overset{\leftarrow}{h}_i = \mathbf{GRU}(e_{x_i}, \overset{\leftarrow}{h}_{i+1}) \quad (2)$$

A brief introduction to NMT

Attention: attention over source/target words:

$$r_{ij} = \mathbf{v}_a^T \tanh(\mathbf{W}_a s_{j-1} + \mathbf{U}_a h_i) \quad (3)$$

$$\alpha_{ij} = \frac{\exp(r_{ij})}{\sum_{i'=1}^{|\mathbf{x}|} \exp(r_{i'j})} \quad (4)$$

Yields “source context vector” c_j at the j -th time step as a weighted sum of all source annotations:

$$c_j = \sum_{i=1}^{|\mathbf{x}|} \alpha_{ij} h_i \quad (5)$$

A brief introduction to NMT

Decoder: another GRU, given the source context vector c_j “unrolls” the target hidden state s_j at time step j :

$$s_j = \mathbf{GRU}(e_{y_{j-1}^*}, s_{j-1}, c_j) \quad (6)$$

Gives probability distribution P_j over all words in the target vocabulary as follows:

$$t_j = g(e_{y_{j-1}^*}, c_j, s_j) \quad (7)$$

$$o_j = \mathbf{W}_o t_j \quad (8)$$

$$P_j = \text{softmax}(o_j) \quad (9)$$

Motivation

NMT models are trained to predict the next word, given the previous context words.

Learn a distribution!

Motivation

However:

At training time: model uses ground truth words as context to predict the next word (context from data distribution)

At inference time: model uses own previous predictions as context (context from model distribution)

This gap is called *exposure bias*.

Motivation

Training typically uses cross-entropy loss, optimizes target sequence to fit ground truth sequence as closely as possible.

Problem: one sentence can have multiple possible translations.

reference: We should comply with the rule.

cand1: We should abide with the rule.

cand2: We should abide by the law.

cand3: We should abide by the rule.

Loss forces prediction back to ground truth (*overcorrection*).

Contributions

Introduce a method that “bridges the gap” between training and inference.

Improves the model’s ability to recover from overcorrection and overall performance.

Proposed method

Simple idea: Instead of just ground truth, feed the model either previous predicted words or ground truth with a certain probability.

Brings training conditions closer to inference time.

Proposed method




Figure: Sample between ground truth word and oracle word

Proposed method

To predict j -th target word y_j :

Recall:

$$s_j = \mathbf{GRU}(e_{y_{j-1}^*}, s_{j-1}, c_j)$$
$$P_j = \text{softmax} \left(\mathbf{W}_o g \left(e_{y_{j-1}^*}, c_j, s_j \right) \right)$$

where s_j next hidden state and P_j probability distribution over target vocabulary

Proposed method

To predict j -th target word y_j :

1. Select an oracle word y_{j-1}^{oracle} at the $\{j-1\}$ -th step.
2. Sample from the ground truth word y_{j-1}^* with a probability of p or from the oracle word y_{j-1}^{oracle} with a probability of $1-p$.
3. Use the sampled word as context $e_{y_{j-1}^*}$.

Oracle word selection

Two strategies to select the oracle words:

1. word-level oracle (greedy search),
2. sentence-level oracle (select an oracle sequence)

Word-level Oracle

Easiest way to pick an oracle word: pick the word with highest probability from the distribution P_{j-1}

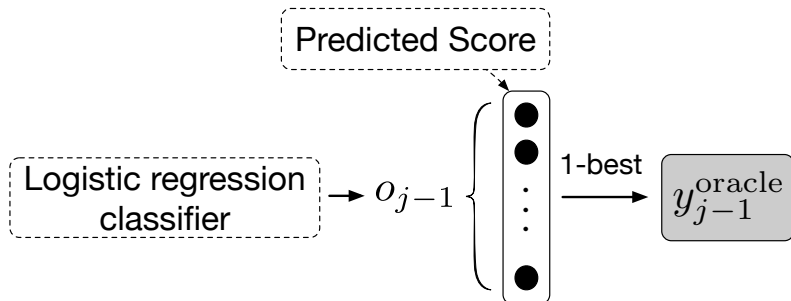


Figure: Word-level oracle without noise

Word-level Oracle

Better: Add noise to the the scores for a better sample




Figure: Word-level oracle with Gumbel noise

Word-level Oracle

Use the *Gumbel-Max* technique: simple, efficient way to sample from categorical distribution

Given Gumbel noise η and a temperature τ , obtain

$$\tilde{o}_{j-1} = (o_{j-1} + \eta) / \tau \quad (10)$$

$$\tilde{P}_{j-1} = \text{softmax}(\tilde{o}_{j-1}) \quad (11)$$

and select the 1-best word from \tilde{P}_{j-1} .

Optimal temperature: $\tau = 0.5$ (found in experiments)

Sentence-level Oracle

Enlarge the search space: perform beam search, apply Gumbel noise at every word generation and get k -best candidate translations

Rank candidates according to some sentence-level metric (here: BLEU)

Force Decoding

d

Sampling with decay

C

Evaluation

a

Conclusion

b

Thank you for your attention.