

Training the face recognition model (s4_train_face_trainer)

Björn Bebensee (2019-21343)

December 3, 2020

0.1 Evaluate the trained network (Week 12) - Step 4

####Designed by Joon Son Chung, November 2020

This is based on https://github.com/joonson/face_trainer. You should read the code if you want to understand more about the training details.

In this step, we train the network on the generated dataset.

The baseline model, available from [here](#), is trained on the VGGFace1 dataset using the softmax loss.

The training and validation sets should also be downloaded to the experiments folder data_dir.

save_path should be changed every time you run a new experiment.

The models take up a significant amount of disk space. Make sure that you have enough space on your Google Drive, and delete any unnecessary/ unsuccessful experiments.

```
[ ]: from google.colab import drive
from zipfile import ZipFile
from tqdm import tqdm
import os, glob, sys, shutil, time
import torch
import torchvision.transforms as transforms
import torchvision.models as models
from PIL import Image
from datetime import datetime # used to timestamp the save_path to keep
    ↳ checkpoints and logs

# mount Google Drive
drive.mount('/content/drive', force_remount=True)

# path of the data directory relative to the home folder of Google Drive
GDRIVE_HOME = '/content/drive/My Drive'
FOLDER      = 'University/2020 Fall/Machine Learning for Visual Understanding/
    ↳ faces'

# specify paths
data_dir    = os.path.join(GDRIVE_HOME, FOLDER) ## path of the general
    ↳ experiment
```

```

initial_model = os.path.join(GDRIVE_HOME, 'University/2020 Fall/Machine Learning_
↳for Visual Understanding/res18_vggface1_baseline.model') ## path to the_
↳pre-trained model
train_zip      = os.path.join(data_dir, 'train.zip') ## training data as zip
val_zip        = os.path.join(data_dir, 'val.zip') ## validation data as zip
now = datetime.now().strftime("%Y%m%d_%H%M%S_") # get the current date and time_
↳as e.g. 20201130_140220_
save_dir = now + 'experiment' # directory name for training logs and checkpoints
save_path     = os.path.join(data_dir, save_dir) ## training logs and trained_
↳models will be saved here

# extract the cropped images
with ZipFile(train_zip, 'r') as zipObj:
    zipObj.extractall("/train_set")
with ZipFile(val_zip, 'r') as zipObj:
    zipObj.extractall("/val_set")
print('Zip extraction complete')

```

Mounted at /content/drive

Zip extraction complete

Make sure that the files have been extracted properly. Make sure that this is a reasonable number.

```

[ ]: train_files = glob.glob('/train_set/*.jpg')
val_files = glob.glob('/val_set/*.jpg')
print(len(train_files), 'train set files and', len(val_files), 'validation set_
↳files found.')

```

58102 train set files and 2431 validation set files found.

First, clone the face recognition trainer from GitHub and add it to path.

```

[ ]: !rm -rf face_trainer
!git clone https://github.com/joonson/face_trainer.git

sys.path.append('face_trainer')

```

Cloning into 'face_trainer'...

remote: Enumerating objects: 30, done.

remote: Counting objects: 100% (30/30), done.

remote: Compressing objects: 100% (23/23), done.

remote: Total 30 (delta 7), reused 25 (delta 5), pack-reused 0

Unpacking objects: 100% (30/30), done.

The training script. Please do not change, but try to read and understand.

```

[ ]: import datetime
from utils import *
from EmbedNet import *

```

```

from DatasetLoader import get_data_loader
import torchvision.transforms as transforms

# ## =====
# ## Trainer script
# ## =====

def train_network(args):

    ## Make folders to save the results and models
    args.model_save_path = args.save_path+"/model"
    args.result_save_path = args.save_path+"/result"

    if not(os.path.exists(args.model_save_path)):
        os.makedirs(args.model_save_path)

    if not(os.path.exists(args.result_save_path)):
        os.makedirs(args.result_save_path)

    ## Load models
    s = EmbedNet(**vars(args)).cuda();

    ## Write args to scorefile
    scorefile = open(args.result_save_path+"/scores.txt", "a+");

    strtime = datetime.datetime.now().strftime("%Y-%m-%d %H:%M:%S")
    scorefile.write('%s\n%s\n'%(strtime,args))
    scorefile.flush()

    ## Input transformations for training
    train_transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Resize(256),
         transforms.RandomCrop([224,224]),
         transforms.RandomRotation(45), # added this transform
         transforms.ColorJitter(brightness=0.2, hue=0.2, saturation=0.2), #
→added this transform
         transforms.RandomGrayscale(p=0.1), # added this transform
         transforms.RandomErasing(p=0.1, scale=(0.02, 0.1), ratio=(0.3, 3.3)), #
→added this transform
         transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
→225]))])

    ## Input transformations for evaluation
    test_transform = transforms.Compose(
        [transforms.ToTensor(),
         transforms.Resize(256),

```

```

        transforms.CenterCrop([224,224]),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
→225]))

    ## Initialise trainer and data loader
    trainLoader = get_data_loader(transform=train_transform, **vars(args));
    trainer      = ModelTrainer(s, **vars(args))

    ## If initial model is specified, start from that model
    if(args.initial_model != ""):
        trainer.loadParameters(args.initial_model);
        print("Model %s loaded!"%args.initial_model);

    besteer    = 100
    bestmodel  = ''

    ## Core training script
    for it in range(1,args.max_epoch+1):

        clr = [x['lr'] for x in trainer.__optimizer__.param_groups]

        print("Training epoch %d with LR %f"%(it,max(clr)));

        loss, traineer = trainer.train_network(trainLoader, verbose=True);

        if it % args.test_interval == 0:

            snapshot = args.model_save_path+"/model%09d.model"%it

            sc, lab = trainer.evaluateFromList(transform=test_transform,
→**vars(args))
            result = tuneThresholdfromScore(sc, lab, [1, 0.1]);

            print("IT %d, VEER %2.4f"%(it, result[1]));
            scorefile.write("IT %d, VEER %2.4f\n"%(it, result[1]));

            trainer.saveParameters(snapshot);

            if result[1] < besteer:
                besteer    = result[1]
                bestmodel  = snapshot

            print("TEER/TAcc %2.2f, TLOSS %f"%( traineer, loss));
            scorefile.write("IT %d, TEER/TAcc %2.2f, TLOSS %f\n"%(it, traineer,
→loss));

            scorefile.flush()

```

```
scorefile.close();

print('Best validation EER: %.4f, %s'%(besteer,bestmodel))
```

Specify the input arguments to the trainer below, and run to train the network. The validation losses will be printed below, and also saved to save_path.

See [here](#) for the meanings of each of these arguments, and see [here](#) for the list of available loss functions. If you use meta-learning loss functions, nPerClass must be 2 or more.

Note that the trainer includes a script to make sure that there are only nPerClass images per class per mini-batch. This helps with the meta-learning loss functions.

```
[ ]: import easydict
args = easydict.EasyDict({ "batch_size": 256, # batch size
                           "trainfunc": "softmaxproto", # loss function
                           "lr": 1e-3, # learning rate
                           "lr_decay": 0.9, # how much to decrease the learning
                           →rate, every 5 epochs
                           "weight_decay": 0, # regularization to reduce
                           →overfitting (e.g. 1e-4 might be reasonable)
                           "margin": 0.2, # for AM-softmax and AAM-softmax
                           "scale": 30, # for AM-softmax and AAM-softmax
                           "nPerClass": 2, # support set + query set size for
                           →meta-learning
                           "nClasses": 1007, #1010, # number of identities in the
                           →training dataset

                           # Don't change below here!!
                           "save_path": save_path,
                           "max_img_per_cls": 500,
                           "nDataLoaderThread": 5,
                           "test_interval": 3,
                           "max_epoch": 15,
                           "optimizer": "adam",
                           "scheduler": "steplr",
                           "hard_prob": 0.5,
                           "hard_rank": 10,
                           "initial_model": initial_model,
                           "train_path": "/train_set",
                           "train_ext": "jpg",
                           "test_path": "/val_set",
                           "test_list": "/val_set/test_list.csv",
                           "model": "ResNet18",
                           "nOut": 512,
                           "mixedprec": False})

train_network(args)
```

Initialised Softmax Loss
 Initialised AngleProto
 Initialised SoftmaxPrototypical Loss
 58102 files from 1007 classes found.
 Initialised Adam optimizer
 Initialised step LR scheduler
 __L__.fc.weight is not in the model.
 __L__.fc.bias is not in the model.
 Model /content/drive/My Drive/University/2020 Fall/Machine Learning for Visual Understanding/res18_vggface1_baseline.model loaded!
 Training epoch 1 with LR 0.001000
 Processing (24320) Loss 8.816987 TEER/TAcc 20.600% - 166.03 Hz
 TEER/TAcc 20.60, TLOSS 8.816987
 Training epoch 2 with LR 0.001000
 Processing (24320) Loss 5.123374 TEER/TAcc 61.589% - 166.11 Hz
 TEER/TAcc 61.59, TLOSS 5.123374
 Training epoch 3 with LR 0.001000
 Processing (24320) Loss 4.172480 TEER/TAcc 73.487% - 166.07 Hz
 Reading 2400 of 2431: 70.43 Hz, embedding size 512
 Computing 46500 of 46589: 17344.58 Hz
 IT 3, VEER 10.8638
 TEER/TAcc 73.49, TLOSS 4.172480
 Training epoch 4 with LR 0.000900
 Processing (24320) Loss 3.712751 TEER/TAcc 78.727% - 147.75 Hz
 TEER/TAcc 78.73, TLOSS 3.712751
 Training epoch 5 with LR 0.000900
 Processing (24320) Loss 3.445094 TEER/TAcc 81.856% - 147.83 Hz
 TEER/TAcc 81.86, TLOSS 3.445094
 Training epoch 6 with LR 0.000900
 Processing (24320) Loss 3.243493 TEER/TAcc 83.657% - 143.87 Hz
 Reading 2400 of 2431: 70.04 Hz, embedding size 512
 Computing 46500 of 46589: 14052.55 Hz
 IT 6, VEER 10.2832
 TEER/TAcc 83.66, TLOSS 3.243493
 Training epoch 7 with LR 0.000810
 Processing (24320) Loss 3.067869 TEER/TAcc 85.504% - 146.53 Hz
 TEER/TAcc 85.50, TLOSS 3.067869
 Training epoch 8 with LR 0.000810
 Processing (24320) Loss 2.933756 TEER/TAcc 86.766% - 145.73 Hz
 TEER/TAcc 86.77, TLOSS 2.933756
 Training epoch 9 with LR 0.000810
 Processing (24320) Loss 2.808039 TEER/TAcc 87.730% - 146.42 Hz
 Reading 2400 of 2431: 69.81 Hz, embedding size 512
 Computing 46500 of 46589: 13626.97 Hz
 IT 9, VEER 9.8344
 TEER/TAcc 87.73, TLOSS 2.808039
 Training epoch 10 with LR 0.000729
 Processing (24320) Loss 2.697170 TEER/TAcc 88.984% - 146.07 Hz

TEER/TAcc 88.98, TLOSS 2.697170
 Training epoch 11 with LR 0.000729
 Processing (24320) Loss 2.592930 TEER/TAcc 89.745% - 146.69 Hz
 TEER/TAcc 89.75, TLOSS 2.592930
 Training epoch 12 with LR 0.000729
 Processing (24320) Loss 2.549067 TEER/TAcc 90.171% - 145.17 Hz
 Reading 2400 of 2431: 71.32 Hz, embedding size 512
 Computing 46500 of 46589: 14418.99 Hz
 IT 12, VEER 9.6728
 TEER/TAcc 90.17, TLOSS 2.549067
 Training epoch 13 with LR 0.000656
 Processing (21760) Loss 2.454362 TEER/TAcc 91.045% - 124.69 Hz

First, evaluate the baseline model yourself. (The baseline model has been trained with this trainer using the softmax loss on the images of the VGGFace dataset cropped to the same specification as our dataset.) Then, try various experiments and record the results. You can try as many or few experiments as you like.

Optional: try to see if you can get better performance by altering the data augmentation. In the current structure this will need to be changed inside `train_network` function.

The parameters in the table denote **changes** from the default parameters that were given with the notebook.

#	Params	Val EER	Notes
0	baseline trained on vggface	22.35	
1	softmax / batch_size=256	11.96	
2	softmax / batch_size=256 / transforms: rotate(45), jitter(0.1), greyscale(0.1), erase(p=0.5,[0.01,0.1])	11.60	
3	softmax / batch_size=200	11.75	
4	softmax / batch_size=100	11.83	
5	softmax / batch_size=100 / weight_decay 1e-4	13.25	Surprisingly perform- ing worse
6	softmax / batch_size=100 / weight_decay=1e-4 / lr=0.0001	12.15	Underfits due to low lr
7	aamsoftmax / batch_size=256 / s=30 m=0.1	10.93	
7	aamsoftmax / batch_size=256 / s=30 m=0.2	11.08	
8	aamsoftmax / batch_size=256 / s=30 m=0.3 / weight_decay=5e-5	10.92	
9	amsoftmax / batch_size=256 / s=30 m=0.1	11.63	
10	angleproto / batch_size=128 / nPerClass=2	10.09	
11	angleproto / batch_size=64 / nPerClass=4	10.53	
12	softmaxproto / batch_size=256 / nPerClass=2	9.82	
13	softmaxproto / batch_size=256 / nPerClass=2 / lr_decay=0.5	9.81	
14	softmaxproto / batch_size=256 / nPerClass=2 / lr_decay=0.5 / weight_decay=5e-4	10.58	

#	Params	Val EER	Notes
15	softmaxproto / batch_size=256 / nPerClass=2 / lr_decay=0.5 / transforms: rotate(45), jitter(0.2), greyscale(0.1), erase(p=0.2,[0.05,0.15])	9.78	Model underfits, try with higher initial lr or lower decay, possibly less erasing (too difficult)
16	softmaxproto / batch_size=256 / nPerClass=2 / lr_decay=0.9 / transforms: rotate(45), jitter(0.2), greyscale(0.1), erase(p=0.1,[0.02,0.1])	9.67	Still underfits, could not finish training in time (score as of epoch 12)
17	ge2e / batch_size=256 / nPerClass=2	10.31	
18	triplet / batch_size=256 / nPerClass=2	10.45	

Choose the best model based on validation performance, and test on test.zip and test_foreign.zip using Step 5 of the code. Write the results below.

Evaluation on the test set:

Baseline model:

test: EER 20.9102

test_foreign: EER 16.4218

Model with best validation performance (#16):

test: EER 8.1117

test_foreign: EER 6.8275

If you have any other thoughts or feedback, please write here.

Although test_foreign is not perfectly “in-distribution”, it seems finetuning on our Korean faces dataset has not only narrowed the gap between the test sets but significantly improved performance on test_foreign as well. Clearly the initial baseline model’s architecture is capable of performing much better on Asian faces indicating that it had been trained on a very biased dataset which includes only very few Asian faces.

Unfortunately I think the time was somewhat constrained so I did not manage to test as extensively with GE2E and Triplet losses as they were somewhat lower on my priority list after seeing results in the AM-Softmax, AAM-Softmax and Angular Prototypical Loss papers. It was not al-

ways possible to run multiple experiments at once due to memory constraints (I think Colab caps this across notebooks?).

I think it could have been interesting to test on the original test set of faces as well (that the baseline was trained on) just to see how much the model is actually forgetting about faces not in the distribution during finetuning.